

Deep Learning Using TensorFlow




Dr. Ash Pahwa

Section 2

Lesson 2.1: Gradient Descent Algorithm

Backpropagation Algorithm:



- Conceptually the backpropagation algorithm is very simple
- Algorithm
 - Assign random values to all the weights of the NN
 - Take the first observed data
 - Forward Propagation: Compute Output
 - Compute error = $(\text{Computed Output} - \text{Observed Output})^2$
 - Backpropagation: adjust weights to reduce error
 - Repeat forward, backward propagation, till error is minimized
 - Repeat the previous step for the next sample till all samples are processed
 - The final weights of the NN will be used for prediction



What is a Gradient?




Gradient Vector: 2 Variables

- Definition: 2 variables x, y
 - The gradient vector of a function $y=f(x)$
 - $\nabla y = \nabla f(x) = \frac{\partial f}{\partial x} = \frac{df}{dx}$
- Gradient property
 - Gradient vector gives the direction of fastest increase (or decrease) of function $f(x)$

Example: Gradient

- $y = x^2 + 6x$
 - $\frac{dy}{dx} = 2x + 6$
 - $x = -4$
 - $y = (-4)^2 + 6(-4) = 16 - 24 = -8$
 - $\frac{dy}{dx} = 2x + 6 = 2 * (-4) + 6 = -2$
 - Gradient at point $(-4, -8) = -2$



What is Gradient Descent Algorithm?



Who Invented Gradient Descent Algorithm?

- Gradient Descent algorithm was invented by **Cauchy** in 1847
- Méthode générale pour la résolution **des** systèmes d'équations simultanées. pp. 536–538

Augustin-Louis Cauchy



Cauchy around 1840. Lithography by Zéphirin Belliard after a painting by Jean Roller.

Born	21 August 1789 Paris, France
Died	23 May 1857 (aged 67) Sceaux, France
Nationality	French
Alma mater	École Nationale des Ponts et Chaussées
Known for	See list
Spouse(s)	Aloise de Bure
Children	Marie Françoise Alicia, Marie Mathilde

What is Gradient Descent Algorithm?

- Suppose a function $y = f(x)$ is given
 - Gradient Descent algorithm allows us to find the values of 'x' where the 'y' value becomes minimum or maximum values
- The Gradient Descent algorithm can be extended to any function with 2 or more variables ' $z = f(x, y)$ '

- Function $y=f(x)$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \parallel_{x^t}$

- Function $z=f(x,y)$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x and y
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial z}{\partial x} \parallel_{x^t, y^t}$
 - $y^{t+1} \leftarrow y^t - \eta \frac{\partial z}{\partial y} \parallel_{x^t, y^t}$



Finding Minimum of a Function Using Gradient Descent Algorithm


2 Variables (x, y) Function

Example 1

Find the value of 'x' where the value of 'y' is minimum

- $y = x^2$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x = 0$
 - $x = 0$
- To find the value of 'x'
 - Where the value of 'y' is minimum
 - Correct answer: $x = 0$

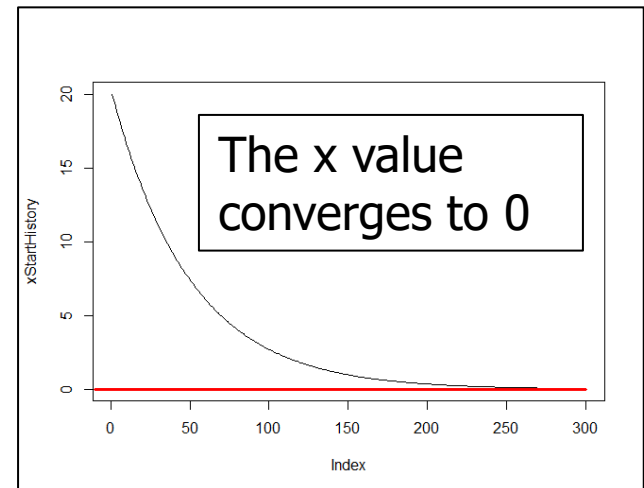

```
> x = seq(-5,5,0.1)
> y = x^2
> dy_dx = function (w1) { 2*w1 }
> plot(x,y,type='l',xlim=c(-10,10),ylim=c(-10,10))
> lines(x,dy_dx(x))
> lines(c(0,0),c(-10,10),lwd=3)
> lines(c(-10,10),c(0,0),lwd=3)
```



Example 1: R Code

- $y = x^2$
- To find minimum point, equate the first derivative = 0
 - $\frac{dy}{dx} = 2x = 0$
 - $x = 0$
- Gradient Descent Algorithm
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \big|_{x^t}$

```
> dy_dx = function (w1) { 2*w1 }  
> xStart = 20  
> learningRate = 0.01  
> maxLimit = 300  
> xStartHistory = rep(0,maxLimit)  
> for ( i in 1:maxLimit )  
+ {  
+   xStartHistory[i] = xStart  
+   xStart = xStart - learningRate * dy_dx(xStart)  
+ }  
> plot(xStartHistory,type='l')  
> lines(c(-10,10),c(0,0),lwd=3,col='red')  
> lines(c(maxLimit,maxLimit),c(0,0),lwd=3,col='red')  
> lines(c(0,maxLimit),c(0,0),lwd=3,col='red')
```





Gradient Descent Algorithm

Parameters

- Parameters
 - Initial value of ' x '
 - Learning Rate
- If the choice of initial value of ' x ' and learning rate is different
 - The Gradient Descent algorithm may not converge



Finding Minimum + Maximum of a Function Using Gradient Descent Algorithm

2 Variables (x,y) Function


Example 3

Find the value of 'x' where the value of 'y' is

* Minimum

* Maximum

- $y = 3x^4 - 16x^3 + 18x^2$
- $\frac{dy}{dx} = 12x^3 - 48x^2 + 36x$
- To find the value of 'x'
 - Where the value of 'y' is local minimum
 - Correct answer: $x = 3, y = -27$




- To find the value of 'x'
 - Where the value of 'y' is local maximum
 - Correct answer: $x = 1, y = 5$

```
> x = seq(0, 4, 0.1)
> y = 3*x^4 - 16*x^3 + 18*x^2
> dy_dx = function (w1) { 12*w1^3 - 48*w1^2 + 36*w1 }
> plot(x, y, type='l', xlim=c(0, 5), ylim=c(-50, 50))
> #lines(x, dy_dx(x))
> lines(c(0, 0), c(-40, 40), lwd=3)
> lines(c(-10, 10), c(0, 0), lwd=3)
> lines(c(3, 3), c(-40, 40), lwd=3, col='red')
> lines(c(1, 1), c(-40, 40), lwd=3, col='red')
> text(1, 42, "Max")
> text(3, 42, "Min")
```


Example 3: Minimum

- $y = 3x^4 - 16x^3 + 18x^2$
- $\frac{dy}{dx} = 12x^3 - 48x^2 + 36x$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial y}{\partial x} \Big|_{x^t}$

```
> dy_dx = function (w1) { 12*w1^3 - 48*w1^2 + 36*w1 }
> xStart = 2
> learningRate = 0.001
> maxLimit = 300
> xStartHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   xStartHistory[i] = xStart
+   xStart = xStart - learningRate * dy_dx(xStart) #
For minimum 'subtract'
+   ## xStart = xStart + learningRate * dy_dx(xStart)
# maximum 'add'
+
+ }
> plot(xStartHistory,type='l')
> lines(c(0,maxLimit),c(3,3),lwd=3,col='red')
```



- To find the value of 'x'
 - Where the value of 'y' is local minimum
 - Correct answer: $x = 3, y = -27$






Finding Minimum of a Function Using Gradient Descent Algorithm

3 Variables (x, y, z) Function

Example 4

Find the value of 'x' where the value of 'y' is
* Minimum

- $z = f(x, y) = x^2 + y^2 - 2x - 6y + 14$
- $\frac{\partial z}{\partial x} = 2x - 2$
- $\frac{\partial z}{\partial y} = 2y - 6$
- To find the value of 'x' and 'y'
 - Where the value of 'z' is local minimum
 - Correct answer: $x = 1, y = 3$
 - $\frac{\partial z}{\partial x} = 2x - 2 = 0; x = 1$
 - $\frac{\partial z}{\partial y} = 2y - 6 = 0; y = 3$





```
> x = seq(0,5,0.1)
> y = seq(0,10,0.1)
> z = function (x,y) { x^2 + y^2 - 2*x - 6*y + 14 }
> dz_dx = function (x1,y1) { 2*x1 - 2 }
> dz_dy = function (x1,y1) { 2*y1 - 6 }
> z<-outer(x,y,z)
> persp(x, y, z)
> contour(z)
```

Example 4: Minimum

- $z = f(x, y) = x^2 + y^2 - 2x - 6y + 14$
- $\frac{\partial z}{\partial x} = 2x - 2; \quad \frac{\partial z}{\partial y} = 2y - 6$
- To find the value of 'x' and 'y'
 - Where the value of 'z' is local minimum
 - Correct answer: $x = 1, y = 3$
 - $\frac{\partial z}{\partial x} = 2x - 2 = 0; x = 1$
 - $\frac{\partial z}{\partial y} = 2y - 6 = 0; y = 3$

```
> dz_dx = function (x1,y1) { 2*x1 - 2 }
> dz_dy = function (x1,y1) { 2*y1 - 6 }
> xStart = 0.1; yStart = 2
> learningRate = 0.01; maxLimit = 500
> xHistory = yHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit)
+ {
+   xHistory[i] = xStart
+   yHistory[i] = yStart
+   dW = dz_dx(xStart,yStart)
+   db = dz_dy(xStart,yStart)
+
+   xStart = xStart - learningRate * dW
+   yStart = yStart - learningRate * db
+ }
> plot(xHistory,type='l')
> plot(yHistory,type='l')
```



- Function $z=f(x,y)$
- Gradient Descent Algorithm: **Minimum**
 - Initialize the value of x and y
 - Learning rate = η
 - While NOT converged:
 - $x^{t+1} \leftarrow x^t - \eta \frac{\partial z}{\partial x} \big|_{x^t, y^t}$
 - $y^{t+1} \leftarrow y^t - \eta \frac{\partial z}{\partial y} \big|_{x^t, y^t}$

©2020 Dr. Ash Pahwa

Solving Regression Problem Using Gradient Descent Algorithm – 2 Variables




- $y = mx + b$
- y is the explanatory variable
- x is the predictor variable
- m is the slope of the line
- b is the intercept

Computing the Regression Line


Compute: Intercept and Slope

- Residual = Observed value – Computed Value
- Suppose regression equation is
 - $y = mx + b$
 - y is the explanatory variable
 - x is the predictor variable
 - m is the slope of the line
 - b is the intercept
- Residual = $y_i - (mx_i + b)$
- Residual² = $(y_i - (mx_i + b))^2$
- Residuals Sum of Squares = (RSS) = $\sum_{i=1}^N (y_i - (mx_i + b))^2$



Residual Sum of Squares

- *Residuals Sum of Squares* = (RSS) = $\sum_{i=1}^N (y_i - (mx_i + b))^2$
- To find the minimum point of this function,
 - we will take the partial derivative of RSS with respect to 'm' and 'b' and set that to zero.
- The RSS is a convex function and it has a minimum point



Partial Derivatives of the RSS w.r.t. Intercept and Slope

- *Residuals Sum of Squares* = (RSS) = $\sum_{i=1}^N (y_i - (mx_i + b))^2$
- To find the minimum point of this function,
 - we will take the partial derivative of RSS with respect to 'm' and 'b' and set that to zero.

$$\begin{aligned} \blacksquare \quad & RSS = \sum_{i=1}^N (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial b} = \sum_{i=1}^N \frac{\partial}{\partial b} (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial b} = -2 \sum_{i=1}^N (y_i - (mx_i + b)) \end{aligned}$$

$$\begin{aligned} \blacksquare \quad & RSS = \sum_{i=1}^N (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial m} = \sum_{i=1}^N \frac{\partial}{\partial m} (y_i - (mx_i + b))^2 \\ \blacksquare \quad & \frac{\partial RSS(m,b)}{\partial m} = -2 \sum_{i=1}^N (y_i - (mx_i + b))x_i \end{aligned}$$

$$\nabla RSS(b, m) = \begin{bmatrix} \frac{\partial RSS(m, b)}{\partial b} \\ \frac{\partial RSS(m, b)}{\partial m} \end{bmatrix} = \begin{bmatrix} -2 \sum_{i=1}^N (y_i - (mx_i + b)) \\ -2 \sum_{i=1}^N (y_i - (mx_i + b))x_i \end{bmatrix} = 0$$

Regression

Closed Form Solution

- To Compute 'm' and 'b'
 - SET GRADIENT = 0

- $$\nabla RSS(b, m) = \begin{bmatrix} \frac{\partial RSS(m, b)}{\partial b} \\ \frac{\partial RSS(m, b)}{\partial m} \end{bmatrix} = \begin{bmatrix} -2 \sum_{i=1}^N (y_i - (mx_i + b)) \\ -2 \sum_{i=1}^N (y_i - (mx_i + b)) x_i \end{bmatrix} = 0$$

- -----

- Top term

- $$b = \left(\frac{\sum y_i}{N} - m \frac{\sum x_i}{N} \right) = \mu_y - m \mu_x$$

- -----

- Bottom term

- $$m = \frac{\sum y_i x_i - \frac{\sum y_i \sum x_i}{N}}{\sum x_i^2 - \frac{\sum x_i \sum x_i}{N}} = r \frac{\sigma_y}{\sigma_x} = \text{Correlation} \frac{\text{Std Dev of } y}{\text{Std Dev of } x}$$

- -----

Closed Form Solution

$$m = \frac{\sum y_i x_i - \frac{\sum y_i \sum x_i}{N}}{\sum x_i^2 - \frac{\sum x_i \sum x_i}{N}}$$

$$b = \left(\frac{\sum y_i}{N} - m \frac{\sum x_i}{N} \right)$$

Regression Equation
 $y = 5x - 1$

	A	B	C	D	E	F	G
1							
2							
3		X	Y		X*Y		X^2
4		0	1		0		0
5		1	3		3		1
6		2	7		14		4
7		3	13		39		9
8		4	21		84		16
9							
10	SUM	10	45		140		30
11	AVERAGE	2	9		28		6
12	StdDev	1.58113883	8.124038405				
13	Correlation	0.97312368					
14							

C16 =E10-(B10*C10/5)							
	A	B	C	D	E	F	G
14							
15	Closed Form	Slope : Using SUM					
16		Numerator	50		(Sum of X*Y) - (1/N)*((Sum of X) * (Sum of Y))		
17		Denominator	10		(Sum of X^2) - (1/N)*((Sum of X * Sum of X))		
18		Slope	5				
19							
20		Intercept	-1		(Mean of Y) - slope * (Mean of X)		
21							

Systems of Linear Equations Multi Variables

- Systems of Linear Equations (Variables = m, Observations = n)

- $y_1 = (b + m_1x_{11} + m_2x_{12} + \cdots + m_mx_{1m}) + e_1$
- $y_2 = (b + m_1x_{21} + m_2x_{22} + \cdots + m_mx_{2m}) + e_2$
- ...
- $y_n = (b + m_1x_{n1} + m_2x_{n2} + \cdots + m_mx_{nm}) + e_n$

- $$Y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{bmatrix} \quad X = \begin{bmatrix} 1 & x_{11} & \dots & x_{1m} \\ 1 & x_{12} & \dots & x_{2m} \\ \dots & \dots & \dots & \dots \\ 1 & x_{1n} & \dots & x_{nm} \end{bmatrix} \quad A = \begin{bmatrix} b \\ m_1 \\ \dots \\ m_m \end{bmatrix} \quad E = \begin{bmatrix} e_1 \\ e_2 \\ \dots \\ e_n \end{bmatrix}$$

- $$A = (X^T X)^{-1} X^T Y$$

Regression

Using R 'lm' command

```
> x = c(0,1,2,3,4)
> y = c(1,3,7,13,21)
> plot(x,y)
> model = lm(y~x)
> summary(model)
```

Call:

```
lm(formula = y ~ x)
```

Residuals:

```
1  2  3  4  5
2 -1 -2 -1  2
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.0000	1.6733	-0.598	0.59220
x	5.0000	0.6831	7.319	0.00527 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 2.16 on 3 degrees of freedom

Multiple R-squared: 0.947, Adjusted R-squared: 0.9293

F-statistic: 53.57 on 1 and 3 DF, p-value: 0.005268

```
> abline(model)
```

```
>
```



Regression Equation



$$y = 5x - 1$$

Regression

Gradient Descent Algorithm Approach

Regression Equation
 $y = 5x - 1$

```
> dRSS_dm = function (m,b) {-2*sum((y-m*x-b)*x) }
> dRSS_db = function (m,b) { -2*sum(y-m*x-b) }
> mStart = bStart = 0
> learningRate = 0.01;  maxLimit = 500
> mHistory = bHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   mHistory[i] = mStart
+   bHistory[i] = bStart
+
+   dW = dRSS_dm(mStart,bStart)
+   db = dRSS_db(mStart,bStart)
+
+   mStart = mStart - learningRate * dW
+   bStart = bStart - learningRate * db
+ }
> plot(mHistory,type='l')
> plot(bHistory,type='l')
```



Solving Regression Problem Using Gradient Descent Algorithm – 2 Variables

Iris Dataset

- $y = mx + b$
 - ***petal.width = m * petal.length + intercept***
- *y is the explanatory variable*
- *x is the predictor variable*
- *m is the slope of the line*
- *b is the intercept*

Read the Iris Dataset

R Code


```
> data(iris)
> #####
> dim(iris)
[1] 150 5
> summary(iris)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
Min. :4.300	Min. :2.000	Min. :1.000	Min. :0.100	setosa :50
1st Qu.:5.100	1st Qu.:2.800	1st Qu.:1.600	1st Qu.:0.300	versicolor:50
Median :5.800	Median :3.000	Median :4.350	Median :1.300	virginica :50
Mean :5.843	Mean :3.057	Mean :3.758	Mean :1.199	
3rd Qu.:6.400	3rd Qu.:3.300	3rd Qu.:5.100	3rd Qu.:1.800	
Max. :7.900	Max. :4.400	Max. :6.900	Max. :2.500	

```
> head(iris)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
> x = iris$Petal.Length
> y = iris$Petal.Width
> plot(x,y)
```



Regression: R

Using R 'lm' command

```
> model = lm(y~x)
> summary(model)


Call:
lm(formula = y ~ x)

Residuals:
    Min       1Q   Median       3Q      Max
-0.56515 -0.12358 -0.01898  0.13288  0.64272

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.363076   0.039762  -9.131  4.7e-16 ***
x             0.415755   0.009582  43.387  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2065 on 148 degrees of freedom
Multiple R-squared:  0.9271, Adjusted R-squared:  0.9266
F-statistic: 1882 on 1 and 148 DF, p-value: < 2.2e-16

> abline(model)
```





*Regression Equation: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*

Regression: R

Gradient Descent Algorithm Approach

```
> dRSS_dm = function (m,b) {-2*sum((y-m*x-b)*x) }
> dRSS_db = function (m,b) { -2*sum(y-m*x-b) }
> mStart = bStart = 0
> learningRate = 0.00001;    maxLimit = 20000
> mHistory = bHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   mHistory[i] = mStart
+   bHistory[i] = bStart
+
+   dW = dRSS_dm(mStart,bStart)
+   db = dRSS_db(mStart,bStart)
+
+   mStart = mStart - learningRate * dW
+   bStart = bStart - learningRate * db
+ }
> plot(mHistory,type='l')
> plot(bHistory,type='l')
> mHistory[maxLimit]
[1] 0.4157522
> bHistory[maxLimit]
[1] -0.3630608
```



*Regression Equation: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*
*Regression Eq: R Grad Desc: Petal.Width = 0.4157 * Petal.Length – 0.3631*

Read the Iris Dataset


Python Code

```
#####  
# Load Libraries  
#  
from sklearn import linear_model  
from sklearn import datasets  
import matplotlib.pyplot as plt  
#####  
# 2. Read the Dataset  
#  
iris = datasets.load_iris()  
  
features = iris["data"]
```

```
petalLength = features[:,2]  
petalLength[0:5]  
Out[13]: array([ 1.4,  1.4,  1.3,  1.5,  1.4])
```

```
petalWidth = features[:,3]  
petalWidth[0:5]  
Out[15]: array([ 0.2,  0.2,  0.2,  0.2,  0.2])
```

```
plt.plot(petalLength,petalWidth,'o')  
Out[16]: [<matplotlib.lines.Line2D at 0x16b604b01d0>]
```



Regression

Using Python 'Scikit-Learn' Library

```
#####  
# 3. Compute the regression equation using Scikit-Learn  
#  
petalLength = petalLength.reshape(-1,1)  
  
petalWidth = petalWidth.reshape(-1,1)  
  
linreg = linear_model.LinearRegression()  
  
linreg.fit(petalLength, petalWidth)  
Out[23]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1,  
normalize=False)  
  
print (linreg.intercept_)  
[-0.36651405]  
  
print (linreg.coef_)  
[[ 0.41641913]]
```

*Regression Equation: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Eq: R Grad Desc: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Equation: Scikit: Petal.Width = 0.4164 * Petal.Length – 0.3665*

Regression: Python

Gradient Descent Algorithm Approach

```
#####
```

```
# 1. Load the libraries
```

```
#
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from sklearn import datasets
```

```
#####
```

```
# 2. Read the Dataset
```

```
#
```

```
iris = datasets.load_iris()
```

```
features = iris["data"]
```

```
x = petalLength = features[:,2]
```

```
x[0:5]
```


```
Out[13]: array([ 1.4,  1.4,  1.3,  1.5,  1.4])
```

```
y = petalWidth = features[:,3]
```

```
y[0:5]
```

```
Out[15]: array([ 0.2,  0.2,  0.2,  0.2,  0.2])
```


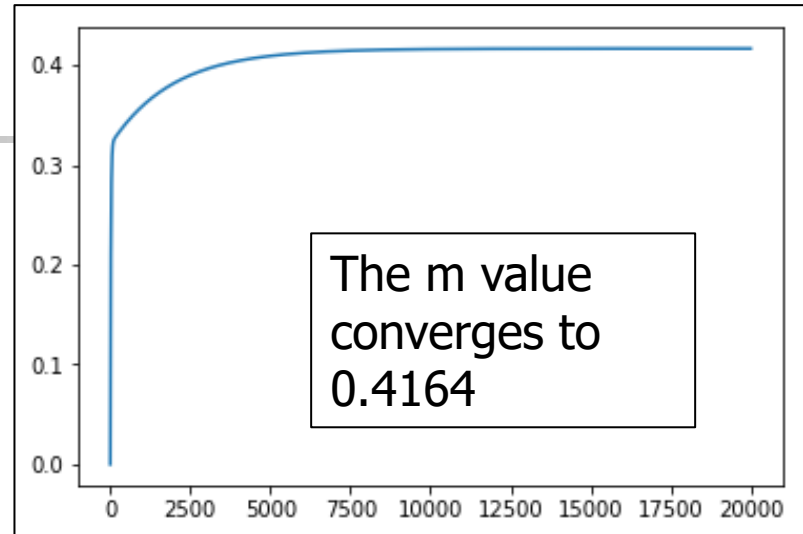
```
plt.plot(x,y,'o')
```



Regression: Python

Gradient Descent Algorithm Approach

```
def dRSS_dm(m,b):  
    return (-2*sum((y-m*x-b)*x))  
  
def dRSS_db(m,b):  
    return (-2*sum((y-m*x-b)))  
  
mStart = 0  
bStart = 0  
learning_rate = 0.00001  
maxLimit = 20000  
mHistory = np.zeros(maxLimit)  
bHistory = np.zeros(maxLimit)  
  
for i in range(maxLimit):  
    mHistory[i] = mStart  
    bHistory[i] = bStart  
    #print(mHistory[i], bHistory[i])  
  
    dW = dRSS_dm(mStart,bStart)  
    db = dRSS_db(mStart,bStart)  
  
    mStart = mStart - learning_rate * dW  
    bStart = bStart - learning_rate * db  
  
print("mHistory=",mHistory[maxLimit-1])  
mHistory= 0.416415833891  
  
print("bHistory=",bHistory[maxLimit-1])  
bHistory= -0.366499084269
```





Final Result

*Regression Eq: R: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Eq: R Grad Desc: Petal.Width = 0.4157 * Petal.Length – 0.3631*

*Regression Eq: Scikit: Petal.Width = 0.4164 * Petal.Length – 0.3665*

*Regression Eq: Python Grad Desc: Petal.Width = 0.4164 * Petal.Length – 0.3665*

Solving Regression Problem Using Gradient Descent Algorithm



Multiple Predictor Variables

- $y = m_1x_1 + m_2x_2 + \dots + m_nx_n + b$
- y is the explanatory variable
- x_1, x_2, \dots, x_n are the predictor variables
- m_1, m_2, \dots, m_n are the coefficients of the predictor variables
- b is the intercept



Compute Residual Sum of Squares

- Residual = Observed value – Computed Value
- Suppose regression equation is
 - $y = m_1x_1 + m_2x_2 + \dots + m_nx_n + b$
 - y is the explanatory variable
 - x_1, x_2, \dots, x_n are the predictor variables
 - m_1, m_2, \dots, m_n are the coefficients of the predictor variables
 - b is the intercept
- $\text{Residual} = y_i - (m_1x_{1i} + m_2x_{2i} + \dots + m_nx_{ni} + b)$
- $\text{Residual}^2 = (y_i - (m_1x_{1i} + m_2x_{2i} + \dots + m_nx_{ni} + b))^2$
- $\text{Residuals Sum of Squares} = (RSS) = \sum_{i=1}^N (y_i - (m_1x_{1i} + m_2x_{2i} + \dots + m_nx_{ni} + b))^2$



Partial Derivatives of the RSS w.r.t. Intercept and Slope

- $RSS = \sum_{i=1}^N (y_i - (m_1 x_{1i} + m_2 x_{2i} + \dots + m_n x_{ni} + b))^2$
- To find the minimum point of this function,
 - we will take the partial derivative of RSS with respect to
 - m_1, m_2, \dots, m_n and
 - 'b' and set that to zero.

- $RSS = \sum_{i=1}^N (y_i - (m_1 x_{1i} + m_2 x_{2i} + \dots + m_n x_{ni} + b))^2$
- $\frac{\partial RSS}{\partial b} = \sum_{i=1}^N \frac{\partial}{\partial b} (y_i - (m_1 x_{1i} + m_2 x_{2i} + \dots + m_n x_{ni} + b))^2$
- $\frac{\partial RSS(m,b)}{\partial b} = -2 \sum_{i=1}^N (y_i - m_1 x_{1i} - m_2 x_{2i} + \dots - m_n x_{ni} - b)$

Partial Derivatives of the RSS w.r.t. Intercept and Slope

- $RSS = \sum_{i=1}^N (y_i - (m_1 x_{1i} + m_2 x_{2i} + \cdots + m_n x_{ni} + b))^2$
- $\frac{\partial RSS}{\partial m_1} = \sum_{i=1}^N \frac{\partial}{\partial m_1} (y_i - (m_1 x_{1i} + m_2 x_{2i} + \cdots + m_n x_{ni} + b))^2$
- -----
- $\frac{\partial RSS}{\partial m_1} = -2 \sum_{i=1}^N (y_i - m_1 x_{1i} - m_2 x_{2i} - \cdots + m_n x_{ni} - b) x_{1i}$
- $\frac{\partial RSS}{\partial m_2} = -2 \sum_{i=1}^N (y_i - m_1 x_{1i} - m_2 x_{2i} - \cdots + m_n x_{ni} - b) x_{2i}$
- ...
- ...
- $\frac{\partial RSS}{\partial m_n} = -2 \sum_{i=1}^N (y_i - m_1 x_{1i} - m_2 x_{2i} - \cdots + m_n x_{ni} - b) x_{ni}$

Data Set

Regression Using 'lm' function

```
> #####  
> x1 = c(0,1,2,3,4)  
> x2 = c(5,10,15,20,15)  
> y = c(1,3,7,13,21)  
> model = lm(y~x1+x2)  
> summary(model)
```

$$\text{Regression Using 'lm' function:}$$
$$y = 6.5x_1 - 0.50x_2 + 2.5$$

Call:

```
lm(formula = y ~ x1 + x2)
```

Residuals:

1	2	3	4	5
1.00e+00	-1.00e+00	-1.00e+00	1.00e+00	-2.22e-16

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.5000	1.9105	1.309	0.321
x1	6.5000	0.8062	8.062	0.015 *
x2	-0.5000	0.2236	-2.236	0.155

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1




Residual standard error: 1.414 on 2 degrees of freedom

Multiple R-squared: 0.9848, Adjusted R-squared: 0.9697

F-statistic: 65 on 2 and 2 DF, p-value: 0.01515

Regression Using Gradient Descent

```
> dRSS_dm1 = function (m1,m2,b) {-2*sum((y - m1*x1 - m2*x2 - b)*x1) }
> dRSS_dm2 = function (m1,m2,b) {-2*sum((y - m1*x1 - m2*x2 - b)*x2) }
> dRSS_db = function (m1,m2,b) {-2*sum (y - m1*x1 - m2*x2 - b) }
> m1Start = m2Start = bStart = -1
> learningRate = 0.0001; maxLimit = 200000
> m1History = m2History = bHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   m1History[i] = m1Start
+   m2History[i] = m2Start
+   bHistory[i] = bStart
+
+   dW1 = dRSS_dm1(m1Start,m2Start,bStart)
+   dW2 = dRSS_dm2(m1Start,m2Start,bStart)
+   db = dRSS_db(m1Start,m2Start,bStart)
+
+   m1Start = m1Start - learningRate * dW1
+   m2Start = m2Start - learningRate * dW2
+   bStart = bStart - learningRate * db
+
+ }
> plot(m1History,type='l')
> plot(m2History,type='l')
> plot(bHistory,type='l')
> m1History[maxLimit]
[1] 6.5
> m2History[maxLimit]
[1] -0.5
> bHistory[maxLimit]
[1] 2.5
```



$$\text{Regression Using Gradient Descent}$$
$$y = 6.5x_1 - 0.50x_2 + 2.5$$



Result

Regression Using 'lm' function: $y = 6.5x_1 - 0.50x_2 + 2.5$

Regression Using Gradient Descent: $y = 6.5x_1 - 0.50x_2 + 2.5$

Solving Regression Problem Using Gradient Descent Algorithm



Multiple Predictor Variables

Iris Dataset

- $y = m_1x_1 + m_2x_2 + \dots + m_nx_n + b$
- y is the explanatory variable
- x_1, x_2, \dots, x_n are the predictor variables
- m_1, m_2, \dots, m_n are the coefficients of the predictor variables
- b is the intercept

Iris Data Set

```
> #####
> # Gradient Decent:
> # 2 predictor variables + 1 response variable
> #
> rm(list=ls(all=TRUE))
> #####
> # Check out the iris dataset
> #
> data(iris)
> #####
> dim(iris)
[1] 150    5
> summary(iris)
  Sepal.Length    Sepal.Width    Petal.Length    Petal.Width      Species
Min.   :4.300    Min.   :2.000    Min.   :1.000    Min.   :0.100    setosa      :50
1st Qu.:5.100    1st Qu.:2.800    1st Qu.:1.600    1st Qu.:0.300    versicolor:50
Median :5.800    Median :3.000    Median :4.350    Median :1.300    virginica  :50
Mean   :5.843    Mean   :3.057    Mean   :3.758    Mean   :1.199
3rd Qu.:6.400    3rd Qu.:3.300    3rd Qu.:5.100    3rd Qu.:1.800
Max.   :7.900    Max.   :4.400    Max.   :6.900    Max.   :2.500
> head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1          3.5          1.4          0.2  setosa
2          4.9          3.0          1.4          0.2  setosa
3          4.7          3.2          1.3          0.2  setosa
4          4.6          3.1          1.5          0.2  setosa
5          5.0          3.6          1.4          0.2  setosa
6          5.4          3.9          1.7          0.4  setosa
```

Regression Using 'lm' function:
Petal.Width = 0.449 Petal.Length – 0.082 Sepal.Length – 0.008

Regression Using 'lm' function

```
> x1 = iris$Petal.Length
> x2 = iris$Sepal.Length
> y = iris$Petal.Width
> model = lm(y~x1+x2)
> summary(model)
```

Call:

```
lm(formula = y ~ x1 + x2)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.60598	-0.12560	-0.02049	0.11616	0.59404

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.008996	0.182097	-0.049	0.9607
x1	0.449376	0.019365	23.205	<2e-16 ***
x2	-0.082218	0.041283	-1.992	0.0483 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2044 on 147 degrees of freedom

Multiple R-squared: 0.929, Adjusted R-squared: 0.9281

F-statistic: 962.1 on 2 and 147 DF, p-value: < 2.2e-16

Regression Using Gradient Descent

```
> dRSS_dm1 = function (m1,m2,b) {-2*sum((y - m1*x1 - m2*x2 - b)*x1) }
> dRSS_dm2 = function (m1,m2,b) {-2*sum((y - m1*x1 - m2*x2 - b)*x2) }
> dRSS_db = function (m1,m2,b) {-2*sum (y - m1*x1 - m2*x2 - b) }
> m1Start = m2Start = bStart = -1
> learningRate = 0.0001;    maxLimit = 200000
> m1History = m2History = bHistory = rep(0,maxLimit)
> for ( i in 1:maxLimit )
+ {
+   m1History[i] = m1Start
+   m2History[i] = m2Start
+   bHistory[i] = bStart
+
+   dW1 = dRSS_dm1(m1Start,m2Start,bStart)
+   dW2 = dRSS_dm2(m1Start,m2Start,bStart)
+   db = dRSS_db(m1Start,m2Start,bStart)
+
+   m1Start = m1Start - learningRate * dW1
+   m2Start = m2Start - learningRate * dW2
+   bStart = bStart - learningRate * db
+
+ }
> plot(m1History,type='l')
> plot(m2History,type='l')
> plot(bHistory,type='l')
> m1History[maxLimit]
[1] 0.4493761
> m2History[maxLimit]
[1] -0.08221782
> bHistory[maxLimit]
[1] -0.008995973
```

Regression Using 'lm' function:

Petal.Width = 0.449 Petal.Length – 0.082 Sepal.Length – 0.008

Regression Using Gradient Descent:

Petal.Width = 0.449 Petal.Length – 0.082 Sepal.Length – 0.008



Other Optimization Algorithms

- Stochastic Gradient Descent
- Momentum
- Nesterov Momentum
- AdaGrad
- RMS Prop
- Adam: Adaptive Momentum Estimation



Summary

- What is a Gradient?
- What is Gradient Descent Algorithm?
- Gradient Descent Algorithm
 - Minimum of a 2-Variable Function
 - Minimum & Maximum of a 2-Variable Function
 - Minimum of 3 Variable Function
 - Solving Regression problem – 2 Variables
 - Solving Regression problem – Iris Dataset (2 Variables)
 - Solving Regression problem – Multiple Variables
 - Solving Regression problem – Iris (Multiple Variables)